

myNotebook

August 18, 2024

1 Notebook: Housing Prices Prediction

Table of Contents: 1. Inspecting data (define X, y columns; basics explorations) 2. Train/test data split 3. Feature Engineering 4. Modeling (including comparing Linear vs Random Forest regressions, building models with cross-validation) 5. Conclusions 6. Next steps

2 1. Inspecting data

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: mycsv = 'data/housing.csv'
df = pd.read_csv(mycsv)
df.head()
```

```
[2]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
[3]: df.shape
```

```
[3]: (20640, 10)
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object

dtypes: float64(9), object(1)
memory usage: 1.6+ MB

3 1.1 Inspecting Data: Define X,Y columns

- X: feature columns = {numeric columns, categorical columns}
- y: target column

```
[5]: columns = list(df.columns)
      print(columns)
```

```
['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms',
'population', 'households', 'median_income', 'median_house_value',
'ocean_proximity']
```

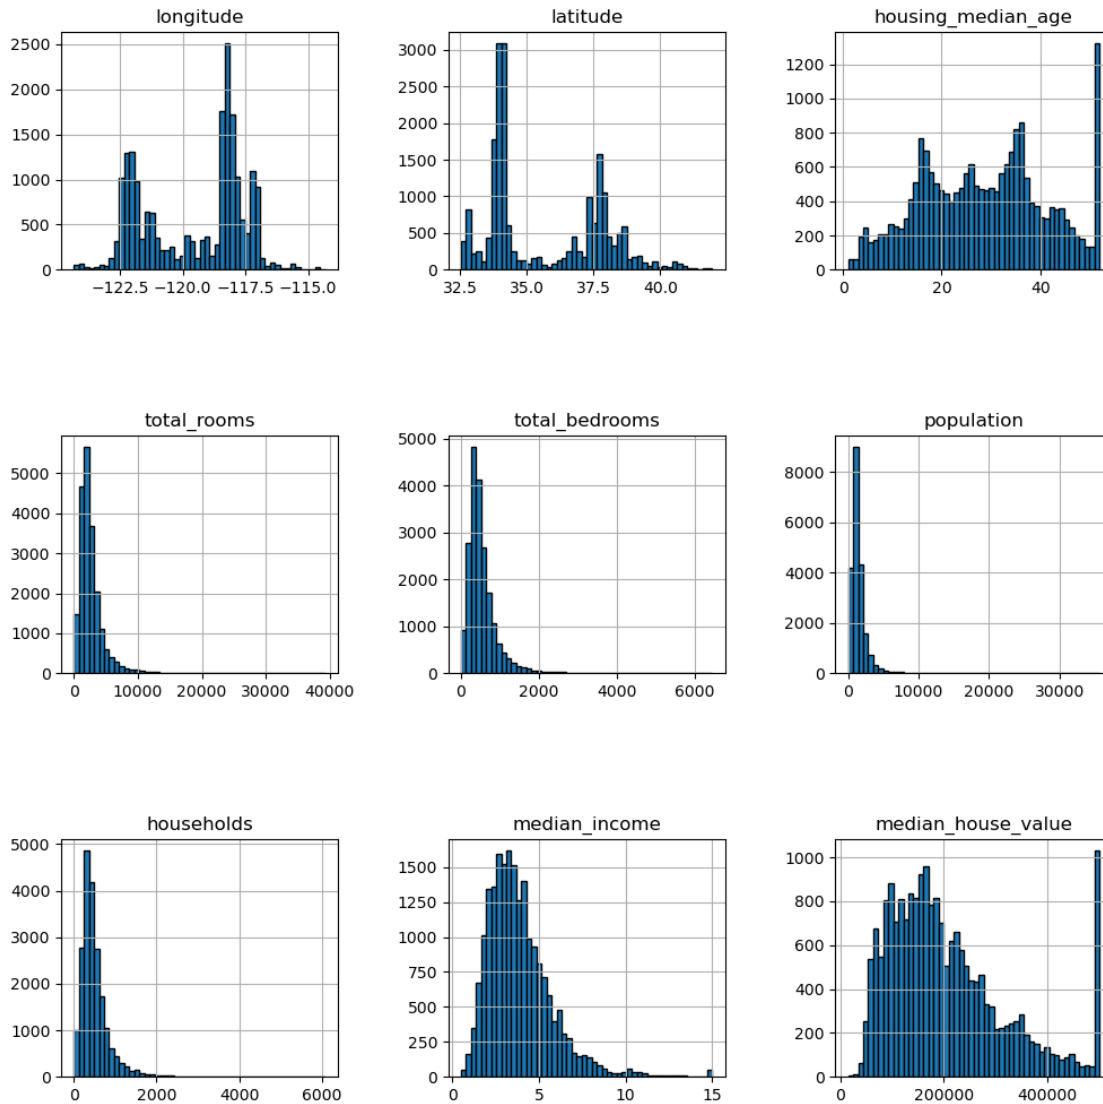
```
[6]: numeric_columns = columns[:-2]
      cat_columns = [columns[-1]]
      y_column = "median_house_value"
      print("Features/num columns:\n", numeric_columns)
      print("Feature/cat columns:\n", cat_columns)
      print("Target column:\n", y_column)
```

```
Features/num columns:
['longitude', 'latitude', 'housing_median_age', 'total_rooms',
'total_bedrooms', 'population', 'households', 'median_income']
Feature/cat columns:
['ocean_proximity']
Target column:
median_house_value
```

4 1.2 Inspecting Data: Basic Exploration

```
[7]: import seaborn as sns
      import matplotlib.pyplot as plt
```

```
[8]: df[numeric_columns + [y_column]].hist(figsize=(12, 12), bins=50,
      ↪ edgecolor="black")
plt.subplots_adjust(hspace=0.7, wspace=0.4)
```



5 1.2.1. Basic Exploration: Linear Correlations between Num Columns and Target Column

```
[9]: correlation_matrix = df[numeric_columns + [y_column]].corr()
correlation_matrix[y_column].sort_values(ascending=False)
```

```
[9]: median_house_value    1.000000
     median_income        0.688075
     total_rooms          0.134153
     housing_median_age    0.105623
     households           0.065843
     total_bedrooms        0.049686
     population           -0.024650
     longitude            -0.045967
     latitude             -0.144160
     Name: median_house_value, dtype: float64
```

6 1.2.2. Basic Exploration: Cat Columns

```
[11]: df[cat_columns].value_counts()
```

```
[11]: ocean_proximity
<1H OCEAN      9136
INLAND         6551
NEAR OCEAN     2658
NEAR BAY       2290
ISLAND          5
     Name: count, dtype: int64
```

7 1.2.3. Basic Exploration: Missing Data. Is it serious?

```
[12]: df.isnull().sum()
```

```
[12]: longitude          0
     latitude           0
     housing_median_age  0
     total_rooms         0
     total_bedrooms     207
     population         0
     households         0
     median_income      0
     median_house_value  0
     ocean_proximity    0
     dtype: int64
```

8 The Heart: from Data to Models to Predictions

Here, we will train models using two different algorithms: * Linear regression with regularization, Lasso * Random forest regression

Then, * Comparing cross-validation (CV) performance of models from the two algorithms * Settling down with the winning algorithm and building model artifacts based on that

9 2. Training/Test data split

```
[13]: from sklearn.model_selection import train_test_split
```

```
[14]: RANDOM_STATE = 123
      TEST_SIZE = 0.2
```

```
[16]: def get_train_test_x_y(df:pd.DataFrame, test_size=TEST_SIZE,
      ↪random_state=RANDOM_STATE):
      train_df, test_df = train_test_split(df, test_size=test_size,
      ↪random_state=random_state)
      x_train = train_df[numeric_columns + cat_columns]
      y_train = train_df[y_column]
      x_test = test_df[numeric_columns + cat_columns]
      y_test = test_df[y_column]
      return x_train, y_train, x_test, y_test
```

```
[17]: x_train, y_train, x_test, y_test = get_train_test_x_y(df)
      print("x_train:", x_train.shape)
      print("y_train:", x_train.shape)
      print("x_test:", x_test.shape)
      print("y_test:", y_test.shape)
```

```
x_train: (16512, 9)
y_train: (16512, 9)
x_test: (4128, 9)
y_test: (4128,)
```

10 3. Feature Engineering

Objective: create a feature-transformer object that transforms raw X_train, X_test

```
[18]: # For num features
      from sklearn.impute import SimpleImputer
      from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      # For cat features
      from sklearn.preprocessing import OneHotEncoder
      from sklearn.compose import ColumnTransformer
```

```
[24]: def get_feature_transformer(x_train:pd.DataFrame):
      num_pipeline = Pipeline([
          ("imputer", SimpleImputer(strategy="median")),
          ("std_scaler", StandardScaler())
      ])

      full_pipeline=ColumnTransformer([
```

```

        ('num', num_pipeline, numeric_columns),
        ('cat', OneHotEncoder(), cat_columns)
    ])

    full_pipeline.fit_transform(x_train)

    return full_pipeline

```

```

[38]: feature_transformer = get_feature_transformer(x_train)
      features = feature_transformer.get_feature_names_out()
      print(f"After transform, there are {len(features)} features:\n", features)

```

After transform, there are 13 features:

```

['num__longitude' 'num__latitude' 'num__housing_median_age'
 'num__total_rooms' 'num__total_bedrooms' 'num__population'
 'num__households' 'num__median_income' 'cat__ocean_proximity_<1H OCEAN'
 'cat__ocean_proximity_INLAND' 'cat__ocean_proximity_ISLAND'
 'cat__ocean_proximity_NEAR BAY' 'cat__ocean_proximity_NEAR OCEAN']

```

```

[39]: transformed_x_train = feature_transformer.transform(x_train)
      print(transformed_x_train.shape)

```

(16512, 13)

11 4. Modeling

```

[41]: from sklearn.metrics import mean_squared_error, mean_absolute_error
      from sklearn.model_selection import cross_val_score

```

12 4.1. Pre-Modeling experiments: which of the two regression algorithms works better, Lasso (regularized Linear Regression) or Random Forest?

12.1 Will do 5-fold cross-validation

```

[42]: CV = 5

```

```

[56]: import warnings
      warnings.filterwarnings('ignore')

```

12.2 Experiment set-up: get cross-validation (CV) metrics from Lasso and Random Forest

```

[48]: from sklearn.linear_model import Lasso
      from sklearn.ensemble import RandomForestRegressor

```

```
[73]: # CV-fold cross-validation
      CV = 5
```

```
[52]: N_EXPERIMENTS = 5
```

```
[53]: RANDOM_STATES = np.random.randint(0, 100, N_EXPERIMENTS)
      RANDOM_STATES
```

```
[53]: array([26, 59, 78, 16, 93])
```

```
[58]: def get_cv_train_rmse_mae(
      model_class:str,
      cv:int=CV,
      n_experiments:int=N_EXPERIMENTS,
      random_states:"List[int] of length n_experiments"=RANDOM_STATES
  ) -> "List[int]":
      if model_class.lower() == "lasso":
          reg_model = Lasso()
      if model_class.lower() == "randomforest":
          reg_model = RandomForestRegressor()

      train_rmse = [0.0] * n_experiments
      train_mae = [0.0] * n_experiments
      test_rmse = [0.0] * n_experiments
      test_mae = [0.0] * n_experiments

      for i in range(n_experiments):
          x_train, y_train, x_test, y_test = get_train_test_x_y(df,
          ↪random_state=random_states[i])

          feature_transformer = get_feature_transformer(x_train)

          transformed_x_train = feature_transformer.transform(x_train)

          reg_scores = cross_val_score(reg_model, transformed_x_train, y_train,
          ↪scoring='neg_mean_squared_error', cv=cv)
          train_rmse[i] = np.sqrt(-reg_scores.mean())

          reg_scores=cross_val_score(reg_model, transformed_x_train, y_train,
          ↪scoring='neg_mean_absolute_error', cv=cv)
          train_mae[i] = -reg_scores.mean()

      return {"train_rmse": train_rmse, "train_mae": train_mae}
```

```
[59]: lasso_cv_metrics = get_cv_train_rmse_mae("lasso")
```

```
[60]: forest_cv_metrics = get_cv_train_rmse_mae("randomforest")
```

```
[64]: print("Lasso regression:")
      print(lasso_cv_metrics)
      print("Random Forest regression:")
      print(forest_cv_metrics)
```

Lasso regression:

```
{'train_rmse': [69335.29106522657, 69244.47427886889, 69085.33095234899,
68609.77538339482, 69068.04777219422], 'train_mae': [50070.4827298464,
50030.33825303199, 49944.77441737708, 49964.73999277492, 50022.066994883615]}
```

Random Forest regression:

```
{'train_rmse': [49377.12748252578, 49380.265901945524, 49586.65237063296,
49778.0672840542, 49819.31350553705], 'train_mae': [32118.778635725128,
32337.593139164826, 32223.577921248838, 32552.66007190974, 32436.120594123728]}
```

12.3 Experiment conclusion: from the results, it seems Random Forest regression outperforms/wins over Lasso

```
[70]: print("Lasso regression:")
      print("RMSE = root mean squared error, MAE = mean absolute error")
      mu, sigma = np.mean(lasso_cv_metrics["train_rmse"]), np.
      ↪std(lasso_cv_metrics["train_rmse"])
      print(f"Train RMSE: {mu:.4f} +/- {sigma:.4f}")
      mu, sigma = np.mean(lasso_cv_metrics["train_mae"]), np.
      ↪std(lasso_cv_metrics["train_mae"])
      print(f"Train MAE: {mu:.4f} +/- {sigma:.4f}")
```

Lasso regression:

RMSE = root mean squared error, MAE = mean absolute error

Train RMSE: 69068.5839 +/- 250.1424

Train MAE: 50006.4805 +/- 45.7347

```
[71]: print("Random Forest regression:")
      print("RMSE = root mean squared error, MAE = mean absolute error")
      mu, sigma = np.mean(forest_cv_metrics["train_rmse"]), np.
      ↪std(forest_cv_metrics["train_rmse"])
      print(f"Train RMSE: {mu:.4f} +/- {sigma:.4f}")
      mu, sigma = np.mean(forest_cv_metrics["train_mae"]), np.
      ↪std(forest_cv_metrics["train_mae"])
      print(f"Train MAE: {mu:.4f} +/- {sigma:.4f}")
```

Random Forest regression:

RMSE = root mean squared error, MAE = mean absolute error

Train RMSE: 49588.2853 +/- 188.2836

Train MAE: 32333.7461 +/- 152.8180

13 4.2. Build Random Forest Regressor with Grid Search Cross-Validation

(since we found Random Forest at least beats Lasso)

with an added training element: use Grid Search Cross Validation to find the optimal model hyperparams

```
[72]: from sklearn.model_selection import GridSearchCV
```

```
[74]: forest_reg = RandomForestRegressor()

param_grid=[
    {'n_estimators': [10, 30, 100], 'max_features':[4, 6, 8, 10]},
    {'bootstrap':[False], 'n_estimators':[3, 10], 'max_features':[2, 3, 4]}
]

grid_search = GridSearchCV(
    forest_reg, param_grid, cv=CV, scoring='neg_mean_absolute_error',
    ↪return_train_score=True
)

grid_search.fit(transformed_x_train, y_train)
```

```
[74]: GridSearchCV(cv=5, estimator=RandomForestRegressor(),
    param_grid=[{'max_features': [4, 6, 8, 10],
    'n_estimators': [10, 30, 100]},
    {'bootstrap': [False], 'max_features': [2, 3, 4],
    'n_estimators': [3, 10]}],
    return_train_score=True, scoring='neg_mean_absolute_error')
```

13.0.1 Just curious: what turned out to be the best hyperparams for Random Forest regressor

```
[75]: grid_search.best_params_
```

```
[75]: {'max_features': 10, 'n_estimators': 100}
```

13.1 4.3. Get the model's test performance

```
[102]: final_model=grid_search.best_estimator_
```

```
[103]: transformed_x_test = feature_transformer.transform(x_test)
y_test_pred = final_model.predict(transformed_x_test)
```

```
[104]: test_rmse = np.sqrt(mean_squared_error(y_true=y_test, y_pred=y_test_pred))
test_mae = mean_absolute_error(y_true=y_test, y_pred=y_test_pred)
```

```
[105]: print("RMSE = root mean squared error, MAE = mean absolute error")
print(f"Test RMSE: {test_rmse:.4f}")
print(f"Test MAE: {test_mae:.4f}")
```

RMSE = root mean squared error, MAE = mean absolute error
Test RMSE: 46857.9324
Test MAE: 30764.8299

13.1.1 Interpretation of Test Performance test_mae = 30765:

On average, the absolute error between predicted and actual house price is \$30765

14 4.4. What can the Random Forest regression model tell us about feature importance?

It is interesting to see: 1. Both linear correlation and Random Forest regressor (a non-linear model) identify median_income as the most important feature 2. Linear correlation gives geo-location (longitude, latitude) weak importance, but Random Forest gives them higher importance.

```
[113]: feature_importances = grid_search.best_estimator_.feature_importances_

cat_encoder = feature_transformer.named_transformers_["cat"]
cat_one_hot_features = list(cat_encoder.categories_[0])
features = numeric_columns + cat_one_hot_features

df_feature_importance = pd.DataFrame(sorted(zip(feature_importances, features),
↪reverse=True))
df_feature_importance.columns = ['Feature', 'Importance']
df_feature_importance
```

```
[113]:
```

	Feature	Importance
0	0.462840	median_income
1	0.155850	INLAND
2	0.109420	longitude
3	0.103742	latitude
4	0.048905	housing_median_age
5	0.033307	population
6	0.026229	total_rooms
7	0.023891	total_bedrooms
8	0.021437	households
9	0.006482	NEAR OCEAN
10	0.005980	<1H OCEAN
11	0.001503	NEAR BAY
12	0.000415	ISLAND

15 5. Conclusions

15.1 Conclusion 0:

For this dataset, Random Forest regression seems outperforms Linear Regression

15.2 Conclusion 1:

The test performance is consistent with, or within the boundary of, previously estimated RMSE/MAE via cross-validation. Recall: we estimated that RMSE: 49588.2853 +/- 188.2836, MAE: 32333.7461 +/- 152.8180

Also, if you compare with the fluctuation of the actual housing prices, the test performance is well within the actual fluctuation. See below, the `std` value

```
[112]: df.describe()[y_column]
```

```
[112]: count      20640.000000
      mean      206855.816909
      std       115395.615874
      min       14999.000000
      25%       119600.000000
      50%       179700.000000
      75%       264725.000000
      max       500001.000000
      Name: median_house_value, dtype: float64
```

15.3 Conclusion 2:

In this machine learning experiment, we have done the following: * Basic data exploration: identify what numeric columns, categorical columns, target column are; if there is severe missing data problem (fortunately, no). * Defined raw data transformation (i.e. feature engineering): scaling numerical features, one-hot encoding categorical features * Conducted experiments to compare which regression algorithms seem to work better: Lasso (regularized linear regression) vs. Random Forest, and settled with Random Forest based on experiment results. * Trained a Random Forest regressor, using cross-validation to choose the optimal hyperparameters. * Found that the test performance (rmse, mae) is within the estimated range.

16 6. Next steps:

- Write model building script to create two main artifacts: (1) feature transformer, (2) ML model - Random Forest regressor
- Write inference script
- Wrap the inference with FastAPI
- Containerize the FastAPI inference app with Docker

```
[ ]:
```